

# Complete AI Web Development Setup Guide

A comprehensive reference for building web applications with AI integration

---

## Table of Contents

1. [VS Code AI Assistant Configuration](#)
  2. [System Prompt \(Complete\)](#)
  3. [Tech Stack Overview](#)
  4. [Free Hosting Options](#)
  5. [Free Database Options](#)
  6. [AI Integration \(Local vs Production\)](#)
  7. [Python AI Client Code](#)
  8. [Flask Route Examples](#)
  9. [Environment Variables](#)
  10. [Project Structure](#)
  11. [Cost Comparison Tables](#)
  12. [MCP Servers Reference](#)
  13. [Quick Reference Commands](#)
- 

## VS Code AI Assistant Configuration

**File:** `new-agent.yaml` or `.vscode/ai-config.yaml`

yaml

*# Free AI Agent Configuration for VS Code*

*# This configuration uses free/open-source AI models*

**name:** Free AI Assistant

**version:** 1.0.0

**schema:** v1

*# System prompt to optimize the agent's behavior*

**systemPrompt:** |

You are an expert programming assistant integrated into VS Code, specializing in full-stack web development.

*## Core Responsibilities:*

- Code writing, debugging, refactoring, and code explanation
- Provide concise yet thorough responses
- Always consider best practices, security implications, and performance

*## Default Tech Stack:*

**\*\*\*Frontend** (unless specified otherwise):\*\*

- HTML5
- SCSS/CSS
- Vanilla JavaScript (or specify framework if needed)
- Responsive design for mobile/tablet/desktop
- React available but may not always be used

**\*\*\*Backend:\*\***

- Python with Flask framework
- SQLAlchemy for database ORM
- Flask-Migrate for database migrations
- RESTful API design

**\*\*\*Database Management:\*\***

- SQLite for local development and prototyping
- PostgreSQL for production (Neon - 10GB free tier)
- MongoDB Atlas (512MB free tier) - alternative for NoSQL
- MySQL via PlanetScale (5GB free tier) - alternative

**\*\*\*AI Integration:\*\***

**\*Development (Local):\***

- Ollama for local AI (free, unlimited, private)
- **Models:** llama3.1, qwen2.5-coder, deepseek-coder
- Runs on localhost:11434

**\*Production (Deployed):\***

- Groq API (FREE tier - primary choice)
  - \* Model: llama-3.1-8b-instant
  - \* Limits: 30 requests/min, 14,400 requests/day
  - \* Cost: \$0 for free tier
- OpenAI API (Paid backup)
  - \* Model: gpt-4o-mini
  - \* Cost: ~\$0.15 per million tokens

**\*Security:\***

- NEVER expose API keys in frontend code
- All AI API calls through backend only
- Use environment variables for sensitive data
- Implement rate limiting (400-500 requests/day per user)

**\*\*Deployment & Hosting:\*\*****\*For Static Landing Pages (FREE):\***

- GitHub Pages, Netlify, or Vercel
- No backend needed
- **Forms:** Formspree, Netlify Forms

**\*For Full-Stack Apps:\***

- Render (free tier) or Railway (\$5 credit)
- Neon PostgreSQL (10GB free)
- Docker for containerization (optional)

**\*For Chat UI with AI:\***

- **Frontend:** GitHub Pages/Netlify
- **Backend:** Flask on Render
- **AI:** Groq (free) or OpenAI
- **Database:** Neon PostgreSQL (optional for history)
- **Rate Limiting:** 400-500 requests/day per user

**\*\*DevOps:\*\***

- Git for version control
- GitHub for repository management
- GitHub Actions for CI/CD
- Docker for containerization (when needed)
- Kubernetes only for large-scale production (usually overkill)

**## Workflow:**

1. **\*\*Always** ask for confirmation\*\* before starting a new task
2. Clarify requirements and scope before coding
3. Explain your approach for complex solutions

4. Provide code with appropriate comments
5. Suggest testing strategies when relevant
6. Prioritize free/low-cost solutions

### ## Best Practices:

- Write clean, maintainable, well-documented code
- Follow security best practices (input validation, secure auth)
- Implement error handling and logging
- Use environment variables for sensitive data
- Never commit secrets to Git
- Consider scalability and performance
- Start simple, scale when needed
- Monitor usage and costs in production

### ## Rate Limiting:

- **Per-user limit:** 400-500 requests per day
- Implement in backend middleware
- Clear error messages when limits exceeded
- Exponential backoff for API retries

*# Models define which AI models this agent can use*

*# These are all free options you can use*

#### models:

*# Option 1: Local models via Ollama (completely free, runs on your machine)*

- **name:** llama-3.2  
**uses:** ollama/llama3.2:3b  
**title:** Llama 3.2 (Fast, Local)
- **name:** qwen-coder  
**uses:** ollama/qwen2.5-coder:7b  
**title:** Qwen Coder (Best for coding)
- **name:** deepseek-coder  
**uses:** ollama/deepseek-coder-v2:16b  
**title:** DeepSeek Coder (Advanced coding)
- **name:** codellama  
**uses:** ollama/codellama:13b  
**title:** Code Llama (Meta's coding model)

*# Option 2: Free API-based models*

- **name:** groq-llama  
**provider:** groq  
**model:** llama-3.1-70b-versatile  
**apiKey:** \${GROQ\_API\_KEY}

```
title: Groq Llama (Fast cloud API)

- name: together-ai
  provider: together
  model: meta-llama/Llama-3.2-11B-Vision-Instruct-Turbo
  apiKey: ${TOGETHER_API_KEY}
  title: Together AI (Free tier available)

- name: openrouter-free
  provider: openrouter
  model: google/gemini-flash-1.5-8b
  apiKey: ${OPENROUTER_API_KEY}
  title: OpenRouter (Very cheap usage)

# Context providers (what the agent can access)
contextProviders:
  - name: codebase
  - name: open-file
  - name: terminal

# Slash commands for quick actions
slashCommands:
  - name: explain
    description: Explain the selected code

  - name: fix
    description: Fix issues in the selected code

  - name: test
    description: Generate tests for the selected code

  - name: refactor
    description: Refactor the selected code

# MCP Servers (optional tools)
mcpServers: []
```

## System Prompt (Complete)

This is the detailed system prompt that guides your AI assistant:

## Core Responsibilities

- Code writing, debugging, refactoring, and code explanation
- Provide concise yet thorough responses
- Always consider best practices, security implications, and performance

## Tech Stack Details

### Frontend

- HTML5
- SCSS/CSS
- Vanilla JavaScript (or specify framework if needed)
- Responsive design for mobile/tablet/desktop
- React available but may not always be used

### Backend

- Python with Flask framework
- SQLAlchemy for database ORM
- Flask-Migrate for database migrations
- RESTful API design
- Async support with asyncio where beneficial

### Database Management

- **SQLite:** Local development and prototyping
- **PostgreSQL:** Production (Neon - 10GB free tier recommended)
- **MongoDB Atlas:** 512MB free tier - alternative for NoSQL
- **MySQL:** PlanetScale (5GB free tier) - alternative

---

## Tech Stack Overview

### Complete Technology Breakdown

Layer	Development	Production
Frontend	HTML/SCSS/JS	Same (static hosting)
Backend	Flask (localhost)	Flask on Render/Railway
Database	SQLite	PostgreSQL (Neon) / MongoDB
AI	Ollama (local)	Groq (free) / OpenAI (paid)
Hosting	localhost:5000	Render/Railway/Netlify
Version Control	Git	GitHub
CI/CD	Manual	GitHub Actions

## Free Hosting Options

### Static Sites (Landing Pages)

Platform	Free Tier	Best For	Limitations
GitHub Pages	Unlimited	Simple static sites	No backend, no server-side code
Netlify	100GB bandwidth/mo	Static + forms	100 form submissions/mo free
Vercel	Unlimited	Next.js, static sites	Serverless functions limited
Cloudflare Pages	Unlimited	Static sites	Learning curve

**Recommended:** GitHub Pages for pure static, Netlify if you need forms

### Full-Stack Apps (with Backend)

Platform	Free Tier	Database	Best For
Render	750 hours/mo	PostgreSQL (90 days)	Flask apps, auto-deploy
Railway	\$5 credit/mo	Included	Small projects, excellent DX
Fly.io	3 VMs free	Extra cost	Python apps, global edge
PythonAnywhere	Limited CPU	MySQL included	Python-specific, beginner-friendly

**Recommended:** Render for production-ready free tier, Railway for best experience

## When to Use What

### Use GitHub Pages when:

- Just a landing page/portfolio
- No backend needed
- No user data to store
- Pure HTML/CSS/JS

**Use Netlify when:**

- Static site + contact forms
- Need some serverless functions
- Want easy deployment

**Use Render when:**

- Need Flask backend
- Need database
- Building chat UI with AI
- Want free production hosting

**Use Railway when:**

- Small project
- Want best developer experience
- \$5/month budget is acceptable

---

## Free Database Options

### Relational Databases

Database	Provider	Free Tier	Best For
SQLite	Local file	Unlimited	Development, prototypes
PostgreSQL	Neon	10GB	Production apps (recommended)
PostgreSQL	Supabase	500MB	Apps needing auth/realtime
PostgreSQL	ElephantSQL	20MB	Tiny projects
PostgreSQL	Render	90 days	Testing (expires)
MySQL	PlanetScale	5GB	Production MySQL needs
MySQL	PythonAnywhere	Included	Simple projects

### NoSQL Databases

Database	Provider	Free Tier	Best For
MongoDB	Atlas	512MB	Document storage, NoSQL
Firebase	Google	1GB + 50K reads/day	Real-time apps, mobile backends
Redis	Upstash	10K commands/day	Caching, sessions



## Recommendations by Use Case

### For Chat UI with AI:

- **Primary:** Neon PostgreSQL (10GB free, serverless, excellent)
- **Alternative:** MongoDB Atlas (512MB, good for JSON-heavy data)

### For Landing Pages:

- **None needed** - use Formspree for forms

### For E-commerce:

- **PostgreSQL** via Neon (structured data, transactions)

### For Real-time Apps:

- **Firebase Firestore** (real-time sync)

---

## AI Integration (Local vs Production)

### Development (Local)

Use **Ollama** - completely free, unlimited, private





```
bash

# Install Ollama
curl -fsSL https://ollama.com/install.sh | sh




# Pull models
ollama pull llama3.1:latest
ollama pull qwen2.5-coder:7b
ollama pull deepseek-coder-v2:16b

# Run Ollama server
ollama serve
```

### Advantages:

-  Free and unlimited
-  Private (data never leaves your machine)
-  No API keys needed
-  Works offline

### Disadvantages:

-  Requires 4-8GB RAM
-  Slower than cloud APIs
-  Only works on your machine

## Production (Deployed)

### Option 1: Groq (FREE - Recommended)

#### Specs:

- **Cost:** FREE
- **Model:** llama-3.1-8b-instant
- **Limits:** 30 req/min, 14,400 req/day
- **Speed:** Very fast (fastest available)

**Sign up:** <https://console.groq.com/>

```
python

from groq import AsyncGroq

client = AsyncGroq(api_key=os.getenv('GROQ_API_KEY'))

response = await client.chat.completions.create(
    model="llama-3.1-8b-instant",
    messages=[{"role": "user", "content": "Hello!"}]
)
```

#### When to use:

- Small to medium projects
- Budget is tight
- Speed is important
- 14,400 requests/day is enough

## Option 2: OpenAI (Paid)

### Specs:

- **Cost:** ~\$0.15 per million tokens
- **Model:** gpt-4o-mini (cheapest)
- **Limits:** None (pay-per-use)
- **Quality:** Better reasoning

**Sign up:** <https://platform.openai.com/>

```
python

from openai import AsyncOpenAI

client = AsyncOpenAI(api_key=os.getenv('OPENAI_API_KEY'))

response = await client.chat.completions.create(
    model="gpt-4o-mini",
    messages=[{"role": "user", "content": "Hello!"}]
)
```

### When to use:

- Need better quality
- Exceed Groq limits
- Budget allows \$5-20/month

## Comparison Table

Feature	Local (Ollama)	Groq (Free)	OpenAI (Paid)
Cost	\$0	\$0	~\$5-20/mo
Speed	Slow-Medium	Very Fast	Fast
Quality	Good	Good	Best
Privacy	100% Private	Cloud	Cloud
Limits	None	14.4K/day	Pay-per-use
Offline	Yes	No	No
Setup	Complex	Easy	Easy

---

## Python AI Client Code

File: `services/ai_client.py`

```
python
```

```
from typing import Optional
import logging
import httpx
import os
from datetime import datetime, timedelta
from openai import AsyncOpenAI
from groq import AsyncGroq, RateLimitError
import asyncio

logger = logging.getLogger(__name__)

# Configuration
ENVIRONMENT = os.getenv('ENVIRONMENT', 'development')
OLLAMA_BASE_URL = os.getenv('OLLAMA_BASE_URL', 'http://localhost:11434')
OLLAMA_MODEL = os.getenv('OLLAMA_MODEL', 'llama3.1:latest')

# Rate limiting configuration
USER_DAILY_LIMIT = int(os.getenv('USER_DAILY_LIMIT', 500))

class RateLimiter:
    """Simple in-memory rate limiter"""

    def __init__(self, limit_per_day: int = USER_DAILY_LIMIT):
        self.limit = limit_per_day
        self.usage = {}

    def check_limit(self, user_id: str) -> tuple[bool, int]:
        """Check if user is within rate limit"""
        now = datetime.now()

        if user_id not in self.usage:
            self.usage[user_id] = {
                'count': 0,
                'reset_time': now + timedelta(days=1)
            }

        user_data = self.usage[user_id]

        if now >= user_data['reset_time']:
            user_data['count'] = 0
            user_data['reset_time'] = now + timedelta(days=1)
```

```
    if user_data['count'] >= self.limit:
        return False, 0

    user_data['count'] += 1
    remaining = self.limit - user_data['count']
    return True, remaining

def get_remaining(self, user_id: str) -> int:
    """Get remaining requests for user"""
    if user_id not in self.usage:
        return self.limit

    user_data = self.usage[user_id]
    now = datetime.now()

    if now >= user_data['reset_time']:
        return self.limit

    return max(0, self.limit - user_data['count'])

class OllamaClient:
    """Local Ollama client for development"""

    def __init__(self, base_url: str = OLLAMA_BASE_URL, model: str = OLLAMA_MODEL):
        self.base_url = base_url
        self.model = model

    async def chat(self, message: str, context: str = "") -> str:
        try:
            async with httpx.AsyncClient(timeout=30.0) as client:
                prompt = f"{context}\n\nUser: {message}\nAssistant:" if context else message
                response = await client.post(
                    f"{self.base_url}/api/generate",
                    json={
                        "model": self.model,
                        "prompt": prompt,
                        "stream": False,
                        "options": {
                            "temperature": 0.7,
                            "num_predict": 500
                        }
                    }
                )
        except Exception as e:
            return f"Error: {e}"
```

```

        if response.status_code == 200:
            result = response.json()
            return result.get("response", "No response from AI")
        else:
            logger.error(f'Ollama error: {response.status_code}')
            return "AI service temporarily unavailable"

    except httpx.ConnectError:
        logger.error("Cannot connect to Ollama")
        return "AI assistant is offline. Please start Ollama with 'ollama serve'"
    except Exception as e:
        logger.error(f'Error calling Ollama: {e}')
        return "I'm having trouble processing your request right now."

class GroqClient:
    """Groq cloud API client (FREE tier)"""

    def __init__(self):
        api_key = os.getenv('GROQ_API_KEY')
        if not api_key:
            raise ValueError("GROQ_API_KEY not set")

        self.client = AsyncGroq(api_key=api_key)
        self.model = "llama-3.1-8b-instant"

    async def chat(self, message: str, context: str = "", max_retries: int = 3) -> str:
        for attempt in range(max_retries):
            try:
                messages = []
                if context:
                    messages.append({"role": "system", "content": context})
                messages.append({"role": "user", "content": message})

                response = await self.client.chat.completions.create(
                    model=self.model,
                    messages=messages,
                    max_tokens=500,
                    temperature=0.7
                )

                return response.choices[0].message.content

            except RateLimitError as e:
                if attempt < max_retries - 1:
                    # Rate limit error, wait and retry
                    await asyncio.sleep(1)

```

```
        wait_time = 60
        logger.warning(f"Groq rate limited, waiting {wait_time}s")
        await asyncio.sleep(wait_time)
    else:
        return "I'm receiving too many requests. Please try again in a minute."

except Exception as e:
    logger.error(f"Groq error: {e}")
    if attempt < max_retries - 1:
        await asyncio.sleep(2 ** attempt)
    else:
        return "I'm having trouble processing your request."

return "Service temporarily unavailable."

class OpenAIClient:
    """OpenAI API client (paid backup)"""

    def __init__(self):
        api_key = os.getenv('OPENAI_API_KEY')
        if not api_key:
            raise ValueError("OPENAI_API_KEY not set")

        self.client = AsyncOpenAI(api_key=api_key)
        self.model = "gpt-4o-mini"

    async def chat(self, message: str, context: str = "") -> str:
        try:
            messages = []
            if context:
                messages.append({"role": "system", "content": context})
            messages.append({"role": "user", "content": message})

            response = await self.client.chat.completions.create(
                model=self.model,
                messages=messages,
                max_tokens=500,
                temperature=0.7
            )

            return response.choices[0].message.content

        except Exception as e:
            logger.error(f"OpenAI error: {e}")
            return "I'm having trouble processing your request "
```



```
        return "I'm having trouble processing your request."

# Global instances
ai_client = None
rate_limiter = RateLimiter(USER_DAILY_LIMIT)

def init_ai():
    """Initialize AI client based on environment"""
    global ai_client

    try:
        if ENVIRONMENT == 'production':
            provider = os.getenv('AI_PROVIDER', 'groq')

            if provider == 'groq':
                ai_client = GroqClient()
                logger.info(f"✓ Groq initialized - Limit: {USER_DAILY_LIMIT} req/day")
            elif provider == 'openai':
                ai_client = OpenAIClient()
                logger.info(f"✓ OpenAI initialized - Limit: {USER_DAILY_LIMIT} req/day")
            else:
                raise ValueError(f"Unknown AI provider: {provider}")
        else:
            ai_client = OllamaClient()
            logger.info(f"✓ Ollama initialized (development)")

    except Exception as e:
        logger.error(f"✗ Error initializing AI: {e}")
        ai_client = None

async def assistant_reply(message: str, user_id: Optional[str] = None) -> dict:
    """Generate AI assistant reply with rate limiting"""
    try:
        if not ai_client:
            init_ai()

        if not ai_client:
            return {
                'success': False,
                'response': "AI assistant is not configured.",
                'remaining': 0
            }
    
```

```
# Check rate limit (production only)
if ENVIRONMENT == 'production' and user_id:
    is_allowed, remaining = rate_limiter.check_limit(user_id)

    if not is_allowed:
        return {
            'success': False,
            'response': f"Daily limit of {USER_DAILY_LIMIT} requests reached. Try tomorrow.",
            'remaining': 0
        }
    else:
        remaining = rate_limiter.get_remaining(user_id) if user_id else USER_DAILY_LIMIT

# AI context
context = """You are a helpful assistant.
Keep responses concise and friendly (under 500 tokens)."""

# Get AI response
response = await ai_client.chat(message, context)

return {
    'success': True,
    'response': response,
    'remaining': remaining
}

except Exception as e:
    logger.error(f"✗ Error in assistant_reply: {e}")
    return {
        'success': False,
        'response': "I'm having trouble right now. Please try again.",
        'remaining': 0
    }

# Initialize on import
init_ai()
```

## Flask Route Examples

**File:** `routes/chat.py`

```
python
```

```
from flask import Blueprint, request, jsonify
from services.ai_client import assistant_reply, rate_limiter, USER_DAILY_LIMIT
import logging

chat_bp = Blueprint('chat', __name__)
logger = logging.getLogger(__name__)

@chat_bp.route('/api/chat', methods=['POST'])
async def chat():
    """Chat endpoint with rate limiting"""
    try:
        data = request.get_json()

        if not data or 'message' not in data:
            return jsonify({
                'success': False,
                'error': 'Message is required'
            }), 400

        message = data['message'].strip()

        if not message:
            return jsonify({
                'success': False,
                'error': 'Message cannot be empty'
            }), 400

        # Get user identifier
        user_id = data.get('user_email') or data.get('session_id') or request.remote_addr

        # Get AI response
        result = await assistant_reply(message, user_id)

        if result['success']:
            return jsonify({
                'success': True,
                'response': result['response'],
                'remaining_requests': result['remaining']
            }), 200
        else:
            status_code = 429 if 'limit' in result['response'].lower() else 500
            return jsonify({
                'success': False,
```

```
        'error': result['response'],
        'remaining_requests': result['remaining']
    }), status_code

except Exception as e:
    logger.error(f"Chat endpoint error: {e}")
    return jsonify({
        'success': False,
        'error': 'Internal server error'
    }), 500

@chat_bp.route('/api/chat/remaining', methods=['GET'])
def get_remaining():
    """Check remaining requests for user"""
    try:
        user_id = request.args.get('user_id') or request.remote_addr
        remaining = rate_limiter.get_remaining(user_id)

        return jsonify({
            'success': True,
            'remaining': remaining,
            'limit': USER_DAILY_LIMIT
        }), 200

    except Exception as e:
        logger.error(f"Error getting remaining: {e}")
        return jsonify({
            'success': False,
            'error': 'Internal server error'
        }), 500
```

**File:** app.py

```
python

from flask import Flask, render_template
from flask_cors import CORS
from routes.chat import chat_bp
import logging

# Configure logging
logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(name)s - %(levelname)s - %(message)s'
)

app = Flask(__name__)
CORS(app)

# Register blueprints
app.register_blueprint(chat_bp)

@app.route('/')
def index():
    return render_template('index.html')

if __name__ == '__main__':
    app.run(debug=True, host='0.0.0.0', port=5000)
```

---

## Environment Variables

### Development (.env)

```
bash

# Environment
ENVIRONMENT=development

# AI Configuration (Local)
AI_PROVIDER=ollama
OLLAMA_BASE_URL=http://localhost:11434
OLLAMA_MODEL=llama3.1:latest

# Rate Limiting
USER_DAILY_LIMIT=500

# Database (Local)
DATABASE_URL=sqlite:///local.db
```

## Production (.env for Render/Railway)

```
bash

# Environment
ENVIRONMENT=production

# AI Configuration (Cloud)
AI_PROVIDER=groq
GROQ_API_KEY=your-groq-api-key-here
# Or for OpenAI:
# AI_PROVIDER=openai
# OPENAI_API_KEY=your-openai-key-here

# Rate Limiting
USER_DAILY_LIMIT=500

# Database (Production)
DATABASE_URL=postgresql://user:pass@host:5432/dbname

# Flask
FLASK_ENV=production
SECRET_KEY=your-secret-key-here
```

## .env.example (Template for team)

```
bash
```

```
# Copy this to .env and fill in your values
```

```
# Environment (development or production)
```

```
ENVIRONMENT=development
```

```
# AI Provider (ollama, groq, or openai)
```

```
AI_PROVIDER=ollama
```

```
# Ollama (for local development)
```

```
OLLAMA_BASE_URL=http://localhost:11434
```

```
OLLAMA_MODEL=llama3.1:latest
```

```
# Groq API (for production - get free key at https://console.groq.com)
```

```
GROQ_API_KEY=
```

```
# OpenAI API (optional backup - get key at https://platform.openai.com)
```

```
OPENAI_API_KEY=
```

```
# Rate Limiting (requests per user per day)
```

```
USER_DAILY_LIMIT=500
```

```
# Database
```

```
DATABASE_URL=sqlite:///local.db
```

```
# Flask
```

```
FLASK_ENV=development
```

```
SECRET_KEY=change-this-in-production
```

---

## Project Structure

```
your-project/
|
├── .vscode/                # VS Code settings
│   └── ai-config.yaml      # VS Code AI assistant config
|
├── app.py                  # Main Flask application
├── requirements.txt         # Python dependencies
├── .env                    # Environment variables (don't commit!)
├── .env.example            # Template for .env
├── .gitignore              # Git ignore file
├── README.md               # Project documentation
|
├── services/               # Business logic
│   ├── __init__.py
│   └── ai_client.py        # AI integration (Ollama/Groq/OpenAI)
|
├── routes/                 # Flask routes/blueprints
│   ├── __init__.py
│   ├── chat.py             # Chat endpoints
│   └── api.py              # Other API endpoints
|
├── models/                 # Database models
│   ├── __init__.py
│   └── database.py         # SQLAlchemy models
|
├── static/                 # Frontend static files
│   ├── css/
│   │   └── styles.css
│   ├── js/
│   │   └── app.js
│   └── images/
|
├── templates/              # HTML templates
│   ├── index.html
│   ├── chat.html
│   └── base.html
|
└── tests/                  # Unit tests
    ├── __init__.py
    ├── test_ai_client.py
    └── test_routes.py
```



## Cost Comparison Tables

### AI API Costs

Provider	Model	Input (per 1M tokens)	Output (per 1M tokens)	Free Tier
Groq	llama-3.1-8b-instant	\$0.05	\$0.08	14.4K req/day FREE
OpenAI	gpt-4o-mini	\$0.15	\$0.60	\$5 credit (expires)
OpenAI	gpt-4o	\$2.50	\$10.00	No free tier
Anthropic	Claude Haiku	\$0.25	\$1.25	No free tier
Local Ollama	Any model	\$0	\$0	Unlimited FREE

### Monthly Cost Examples (Chat UI)

#### Scenario 1: Small Project (100 users/day)

- 100 users × 5 messages = 500 messages/day = 15,000/month
- Average 200 tokens per message = 3M tokens/month

Provider	Monthly Cost
Groq	\$0 (within free tier)
OpenAI (gpt-4o-mini)	~\$0.45
Local Ollama	\$0

#### Scenario 2: Medium Project (1,000 users/day)

- 1,000 users × 5 messages = 5,000 messages/day = 150,000/month
- Average 200 tokens per message = 30M tokens/month

Provider	Monthly Cost
Groq	~\$1.50 (slightly over free tier)
OpenAI (gpt-4o-mini)	~\$4.50
Local Ollama	\$0 (but needs powerful server)

#### Scenario 3: Large Project (5,000 users/day)

- 5,000 users × 5 messages = 25,000 messages/day = 750,000/month
- Average 200 tokens per message = 150M tokens/month

Provider	Monthly Cost
Groq	~\$7.50
OpenAI (gpt-4o-mini)	~\$22.50

## Hosting Costs

Platform	Free Tier	After Free Tier
GitHub Pages	Unlimited	Always free
Netlify	100GB/mo	\$19/mo (Pro)
Vercel	100GB/mo	\$20/mo (Pro)
Render	750 hours/mo	\$7/mo per service
Railway	\$5 credit/mo	Pay-as-you-go
Fly.io	3 VMs free	\$1.94/mo per VM

## Database Costs

Provider	Free Tier	After Free Tier
Neon PostgreSQL	10GB	\$19/mo (Launch)
Supabase	500MB	\$25/mo (Pro)
MongoDB Atlas	512MB	\$9/mo (M2)
PlanetScale	5GB	\$29/mo (Scaler)
SQLite	Unlimited	Always free (local file)

## Total Monthly Costs by Setup

### Setup 1: Landing Page Only

- Hosting: GitHub Pages (free)
- Forms: Formspree (free)
- **Total: \$0/month**

### Setup 2: Chat UI (Small - 100 users/day)

- Frontend: Netlify (free)
- Backend: Render (free)
- Database: Neon PostgreSQL (free)
- AI: Groq (free)
- **Total: \$0/month**

### Setup 3: Chat UI (Medium - 1,000 users/day)

- Frontend: Netlify (free)
- Backend: Render (free or \$7/mo)
- Database: Neon PostgreSQL (free)
- AI: Groq (\$1.50/mo)
- **Total: \$1.50-8.50/month**

### Setup 4: Production App (5,000 users/day)

- Frontend: Vercel (free)
- Backend: Render (\$7/mo)
- Database: Neon PostgreSQL (free or \$19/mo if >10GB)
- AI: Groq (\$7.50/mo)
- **Total: \$14.50-33.50/month**

---

## MCP Servers Reference

### What Are MCP Servers?

**MCP = Model Context Protocol**

A standard created by Anthropic that allows AI assistants to connect to external tools and data sources. Think of them as plugins or extensions.

### Popular MCP Servers

MCP Server	What It Does	Use Case
memory-mcp	Long-term memory across sessions	Remember project context
filesystem-mcp	Read/write files	Auto-update config files
github-mcp	GitHub API access	Create issues, PRs automatically
postgres-mcp	PostgreSQL queries	Query your database
brave-search-mcp	Web search	Research while coding
puppeteer-mcp	Browser automation	Test websites
slack-mcp	Slack integration	Send notifications
google-drive-mcp	Google Drive access	Read/write Drive files

### How to Add MCP Servers

### Example 1: Add Memory

```
yaml

mcpServers:
  - uses: anthropic/memory-mcp
    config:
      max_memory_mb: 100
```

### Example 2: Add GitHub Integration

```
yaml

mcpServers:
  - uses: github-mcp
    config:
      token: ${GITHUB_TOKEN}
      repos:
        - username/repo-name
```

### Example 3: Multiple MCP Servers

```
yaml

mcpServers:
  - uses: anthropic/memory-mcp
  - uses: filesystem-mcp
    config:
      allowed_directories:
        - ~/projects
  - uses: github-mcp
    config:
      token: ${GITHUB_TOKEN}
```

### Should You Use MCP Servers?

#### NO (Keep empty) if:

- You're just starting out
- Basic coding assistance is enough
- You don't need persistent memory
- You want to keep things simple

**YES if:**

- You want AI to remember past conversations
- You want AI to interact with tools automatically
- You're building complex workflows
- Basic functionality isn't enough

**Recommendation:** Start with `mcpServers: []` (empty), add them later when needed

---

## Quick Reference Commands

### Ollama Commands

```
bash

# Install Ollama (Mac/Linux)
curl -fsSL https://ollama.com/install.sh | sh

# Install Ollama (Windows)
# Download from: https://ollama.com/download

# Pull models
ollama pull llama3.1:latest
ollama pull qwen2.5-coder:7b
ollama pull deepseek-coder-v2:16b
ollama pull codellama:13b

# List installed models
ollama list

# Run Ollama server
ollama serve

# Test a model
ollama run llama3.1

# Remove a model
ollama rm llama3.1
```

### Python/Flask Commands

```
bash
```

```
# Create virtual environment
```

```
python -m venv venv
```

```
# Activate virtual environment (Mac/Linux)
```

```
source venv/bin/activate
```

```
# Activate virtual environment (Windows)
```

```
venv\Scripts\activate
```

```
# Install dependencies
```

```
pip install -r requirements.txt
```

```
# Create requirements.txt
```

```
pip freeze > requirements.txt
```

```
# Run Flask app
```

```
python app.py
```

```
# Run Flask app with auto-reload
```

```
flask run --reload
```

```
# Run tests
```

```
pytest
```

## Git Commands

```
bash

# Initialize repository
git init

# Add all files
git add .

# Commit changes
git commit -m "Initial commit"

# Add remote repository
git remote add origin https://github.com/username/repo.git

# Push to GitHub
git push -u origin main

# Create new branch
git checkout -b feature-name

# Switch branches
git checkout main

# Pull latest changes
git pull origin main

# Check status
git status
```

## Docker Commands (Optional)

```
bash

# Build Docker image
docker build -t my-flask-app .

# Run Docker container
docker run -p 5000:5000 my-flask-app

# Run with environment variables
docker run -p 5000:5000 --env-file .env my-flask-app

# Stop container
docker stop container-id

# List running containers
docker ps

# List all containers
docker ps -a

# Remove container
docker rm container-id
```

## Database Commands

```
bash

# SQLite
sqlite3 local.db

# Create tables (using Flask-Migrate)
flask db init
flask db migrate -m "Initial migration"
flask db upgrade

# PostgreSQL (psql)
psql -U username -d database_name

# MongoDB (mongosh)
mongosh "mongodb://localhost:27017"
```

---

## Python Requirements.txt



**File:** requirements.txt

```
# Flask Framework
Flask==3.0.0
Flask-CORS==4.0.0
Flask-Migrate==4.0.5
Flask-SQLAlchemy==3.1.1

# Database
SQLAlchemy==2.0.23
psycopg2-binary==2.9.9 # PostgreSQL
pymongo==4.6.1 # MongoDB (optional)

# AI Libraries
openai==1.6.1 # For OpenAI API
groq==0.4.1 # For Groq API
httpx==0.25.2 # For Ollama

# Utilities
python-dotenv==1.0.0
requests==2.31.0

# Development
pytest==7.4.3
black==23.12.1
flake8==6.1.0
```

**Install with:**

```
bash

pip install -r requirements.txt
```

---

## Example Frontend (Chat UI)

**File:** static/js/app.js

javascript

*// Chat UI JavaScript*

```
class ChatApp {
  constructor() {
    this.chatForm = document.getElementById('chat-form');
    this.chatInput = document.getElementById('chat-input');
    this.chatMessages = document.getElementById('chat-messages');
    this.remainingEl = document.getElementById('remaining-requests');

    this.init();
  }

  init() {
    this.chatForm.addEventListener('submit', (e) => this.handleSubmit(e));
    this.updateRemaining();
  }

  async handleSubmit(e) {
    e.preventDefault();

    const message = this.chatInput.value.trim();
    if (!message) return;

    // Add user message to UI
    this.addMessage(message, 'user');
    this.chatInput.value = '';

    // Show loading
    const loadingId = this.addMessage('Thinking...', 'assistant', true);

    try {
      const response = await fetch('/api/chat', {
        method: 'POST',
        headers: {
          'Content-Type': 'application/json'
        },
        body: JSON.stringify({
          message: message,
          session_id: this.getSessionId()
        })
      });

      const data = await response.json();
```

```
// Remove loading message
document.getElementById/loadingId).remove();

if (data.success) {
  this.addMessage(data.response, 'assistant');
  this.updateRemaining(data.remaining_requests);
} else {
  this.addMessage(`Error: ${data.error}`, 'error');
}

} catch (error) {
  document.getElementById/loadingId).remove();
  this.addMessage('Failed to get response. Please try again.', 'error');
  console.error('Chat error:', error);
}
}

addMessage(text, sender, isLoading = false) {
  const messageId = `msg-${Date.now()}`;
  const messageEl = document.createElement('div');
  messageEl.id = messageId;
  messageEl.className = `message message-${sender}`;
  messageEl.textContent = text;

  if (isLoading) {
    messageEl.classList.add('loading');
  }

  this.chatMessages.appendChild(messageEl);
  this.chatMessages.scrollTop = this.chatMessages.scrollHeight;

  return messageId;
}

async updateRemaining(count = null) {
  if (count !== null) {
    this.remainingEl.textContent = `Remaining: ${count}`;
    return;
  }

  try {
    const response = await fetch(`/api/chat/remaining?user_id=${this.getSessionId()}`);
    const data = await response.json();

    if (data.success) {

```

```
        this.remainingEl.textContent = `Remaining: ${data.remaining}/${data.limit}`;
    }
} catch (error) {
    console.error('Failed to get remaining requests:', error);
}
}

getSessionId() {
    let sessionId = localStorage.getItem('session_id');
    if (!sessionId) {
        sessionId = `session-${Date.now()}-${Math.random()}`;
        localStorage.setItem('session_id', sessionId);
    }
    return sessionId;
}

// Initialize app when DOM is ready
document.addEventListener('DOMContentLoaded', () => {
    new ChatApp();
});
```

**File:** templates/index.html

html

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>AI Chat Assistant</title>
  <link rel="stylesheet" href="{{ url_for('static', filename='css/styles.css') }}">
</head>
<body>
  <div class="container">
    <header>
      <h1>AI Chat Assistant</h1>
      <p id="remaining-requests">Loading...</p>
    </header>

    <div id="chat-messages" class="chat-messages">
      <!-- Messages appear here -->
    </div>

    <form id="chat-form" class="chat-form">
      <input
        type="text"
        id="chat-input"
        placeholder="Type your message..."
        autocomplete="off"
        required
      >
      <button type="submit">Send</button>
    </form>
  </div>

  <script src="{{ url_for('static', filename='js/app.js') }}"></script>
</body>
</html>
```

**File:** static/css/styles.css

## CSS

```
* {
  margin: 0;
  padding: 0;
  box-sizing: border-box;
}

body {
  font-family: -apple-system, BlinkMacSystemFont, 'Segoe UI', Roboto, sans-serif;
  background: linear-gradient(135deg, #667eea 0%, #764ba2 100%);
  min-height: 100vh;
  display: flex;
  justify-content: center;
  align-items: center;
  padding: 20px;
}

.container {
  width: 100%;
  max-width: 800px;
  height: 600px;
  background: white;
  border-radius: 20px;
  box-shadow: 0 20px 60px rgba(0, 0, 0, 0.3);
  display: flex;
  flex-direction: column;
}

header {
  padding: 20px;
  border-bottom: 2px solid #f0f0f0;
  display: flex;
  justify-content: space-between;
  align-items: center;
}

h1 {
  font-size: 1.5em;
  color: #333;
}

#remaining-requests {
  color: #666;
  font-size: 0.9em;
```

```
}

.chat-messages {
  flex: 1;
  padding: 20px;
  overflow-y: auto;
  display: flex;
  flex-direction: column;
  gap: 15px;
}

.message {
  max-width: 70%;
  padding: 12px 16px;
  border-radius: 12px;
  word-wrap: break-word;
}

.message-user {
  align-self: flex-end;
  background: #667eea;
  color: white;
}

.message-assistant {
  align-self: flex-start;
  background: #f0f0f0;
  color: #333;
}

.message-error {
  align-self: center;
  background: #ff6b6b;
  color: white;
}

.message.loading {
  opacity: 0.6;
  animation: pulse 1.5s infinite;
}

@keyframes pulse {
  0%, 100% { opacity: 0.6; }
  50% { opacity: 1; }
}
```

```
.chat-form {
  padding: 20px;
  border-top: 2px solid #f0f0f0;
  display: flex;
  gap: 10px;
}

#chat-input {
  flex: 1;
  padding: 12px 16px;
  border: 2px solid #e0e0e0;
  border-radius: 12px;
  font-size: 1em;
  outline: none;
  transition: border-color 0.3s;
}

#chat-input:focus {
  border-color: #667eea;
}

button {
  padding: 12px 30px;
  background: #667eea;
  color: white;
  border: none;
  border-radius: 12px;
  font-size: 1em;
  font-weight: 600;
  cursor: pointer;
  transition: background 0.3s;
}

button:hover {
  background: #5568d3;
}

button:active {
  transform: scale(0.98);
}

/* Mobile responsive */
@media (max-width: 768px) {
  .container {
    padding: 10px;
  }
}
```



```
height: 100vh;  
border-radius: 0;  
}  
  
.message {  
  max-width: 85%;  
}  
}
```

## .gitignore File

File: `.gitignore`

```
# Python
__pycache__/
*.py[cod]
*$py.class
*.so
.Python
venv/
env/
ENV/
build/
develop-eggs/
dist/
downloads/
eggs/
.eggs/
lib/
lib64/
parts/
sdist/
var/
wheels/
*.egg-info/
.installed.cfg
*.egg

# Environment variables
.env
.env.local
.env.*.local

# IDE
.vscode/
.idea/
*.swp
*.swo
*~

# Database
*.db
*.sqlite
*.sqlite3
local.db

# Logs
*.log
```

```
.log
logs/

# OS
.DS_Store
Thumbs.db

# Flask
instance/
.webassets-cache

# Testing
.pytest_cache/
.coverage
htmlcov/

# Docker
docker-compose.override.yml

# Node (if using)
node_modules/
npm-debug.log
yarn-error.log
```

## Deployment Guide

### Deploy to Render (FREE)

#### Step 1: Prepare Your Code

```
bash

# Make sure you have these files:
# - app.py
# - requirements.txt
# - .env.example (don't include .env!)
```

#### Step 2: Create Render Account

- Go to <https://render.com>
- Sign up with GitHub

#### Step 3: Create New Web Service

- Click "New +" → "Web Service"
- Connect your GitHub repository
- Configure:
  - **Name:** your-app-name
  - **Environment:** Python 3
  - **Build Command:** `pip install -r requirements.txt`
  - **Start Command:** `gunicorn app:app`
  - **Plan:** Free

**Step 4: Add Environment Variables** In Render dashboard, add:

```
ENVIRONMENT=production
AI_PROVIDER=groq
GROQ_API_KEY=your-key-here
USER_DAILY_LIMIT=500
DATABASE_URL=your-neon-url-here
SECRET_KEY=generate-random-string
```

### Step 5: Deploy

- Click "Create Web Service"
- Wait for deployment (5-10 minutes)
- Your app will be live at `https://your-app-name.onrender.com`

## Deploy to Railway (\$5 credit/month)

### Step 1: Create Railway Account

- Go to <https://railway.app>
- Sign up with GitHub

### Step 2: Create New Project

- Click "New Project"
- Select "Deploy from GitHub repo"
- Choose your repository

### Step 3: Add Environment Variables

- Go to "Variables" tab
- Add all environment variables

### Step 4: Deploy

- Railway auto-deploys on push to main branch
- Get your URL from dashboard

## Deploy Frontend to Netlify (FREE)

### For Static Frontend Only:

#### Step 1: Create Netlify Account

- Go to <https://netlify.com>
- Sign up with GitHub

#### Step 2: Deploy

- Drag and drop your `static/` folder
- Or connect GitHub repository

#### Step 3: Configure

- Set build command if needed
- Add environment variables for API URLs

---

## Troubleshooting

### Common Issues

#### Issue: "Cannot connect to Ollama"

```
bash

# Solution: Start Ollama server
ollama serve

# Or check if it's running
ps aux | grep ollama
```

#### Issue: "GROQ\_API\_KEY not set"

```
bash
```

```
# Solution: Add to .env file
```

```
echo "GROQ_API_KEY=your-key-here" >> .env
```

```
# Or export in terminal
```

```
export GROQ_API_KEY=your-key-here
```

### Issue: "Rate limit exceeded"

```
python
```

```
# Solution: Wait 1 minute or increase USER_DAILY_LIMIT
```

```
USER_DAILY_LIMIT=1000 # in .env
```

### Issue: "Database connection failed"

```
bash
```

```
# Solution: Check DATABASE_URL format
```

```
# PostgreSQL: postgresql://user:pass@host:5432/dbname
```

```
# SQLite: sqlite:///local.db
```

### Issue: "Module not found"

```
bash
```

```
# Solution: Install missing package
```

```
pip install package-name
```

```
# Or reinstall all
```

```
pip install -r requirements.txt
```

### Issue: "Port already in use"

```
bash
```

```
# Solution: Kill process on port 5000
```

```
# Mac/Linux:
```

```
lsof -ti:5000 | xargs kill -9
```

```
# Windows:
```






```
netstat -ano | findstr :5000
```

```
taskkill /PID <PID> /F
```





---

## Security Best Practices





### API Keys

-  **DO:** Store in environment variables
-  **DO:** Use `.env` file (add to `.gitignore`)
-  **DON'T:** Commit to Git
-  **DON'T:** Hardcode in source files
-  **DON'T:** Expose in frontend code




### Rate Limiting

-  **DO:** Implement per-user limits
-  **DO:** Set reasonable daily limits (400-500)
-  **DO:** Log violations
-  **DON'T:** Allow unlimited requests





### Input Validation

-  **DO:** Validate all user inputs
-  **DO:** Sanitize before processing
-  **DO:** Set max message length
-  **DON'T:** Trust user input blindly

### HTTPS

-  **DO:** Use HTTPS in production
-  **DO:** Let platforms handle SSL (Render/Netlify)
-  **DON'T:** Send sensitive data over HTTP

### Database

-  **DO:** Use parameterized queries
  -  **DO:** Hash passwords (bcrypt)
  -  **DO:** Regular backups
  -  **DON'T:** Store plain-text passwords
-

## Performance Optimization

### Caching

```
python

from flask_caching import Cache

cache = Cache(app, config={'CACHE_TYPE': 'simple'})

@cache.cached(timeout=300) # Cache for 5 minutes
def get_popular_responses():
    # Expensive operation
    pass
```

### Database Indexing

```
python

# Add indexes to frequently queried fields
class User(db.Model):
    email = db.Column(db.String(120), index=True, unique=True)
```

### Async Operations

```
python

# Use async for I/O-bound operations
async def fetch_multiple_apis():
    tasks = [
        fetch_api_1(),
        fetch_api_2(),
        fetch_api_3()
    ]
    results = await asyncio.gather(*tasks)
    return results
```

---

## Testing

### Unit Tests Example

**File:** `tests/test_ai_client.py`



```
python

import pytest
from services.ai_client import RateLimiter, assistant_reply

def test_rate_limiter():
    limiter = RateLimiter(limit_per_day=2)

    # First request - should succeed
    allowed, remaining = limiter.check_limit("user1")
    assert allowed == True
    assert remaining == 1

    # Second request - should succeed
    allowed, remaining = limiter.check_limit("user1")
    assert allowed == True
    assert remaining == 0

    # Third request - should fail
    allowed, remaining = limiter.check_limit("user1")
    assert allowed == False
    assert remaining == 0

@pytest.mark.asyncio
async def test_assistant_reply():
    response = await assistant_reply("Hello", user_id="test_user")

    assert response['success'] == True
    assert 'response' in response
    assert 'remaining' in response
```

### Run tests:

```
bash

pytest
pytest -v # Verbose
pytest tests/test_ai_client.py # Specific file
```

---

## Resources & Links

### Documentation

- **Flask:** <https://flask.palletsprojects.com/>
- **SQLAlchemy:** <https://www.sqlalchemy.org/>
- **Ollama:** <https://ollama.com/>
- **Groq:** <https://console.groq.com/docs>
- **OpenAI:** <https://platform.openai.com/docs>

## Free Hosting

- **Render:** <https://render.com>
- **Railway:** <https://railway.app>
- **Netlify:** <https://netlify.com>
- **Vercel:** <https://vercel.com>
- **GitHub Pages:** <https://pages.github.com>

## Free Databases

- **Neon PostgreSQL:** <https://neon.tech>
- **Supabase:** <https://supabase.com>
- **MongoDB Atlas:** <https://www.mongodb.com/atlas>
- **PlanetScale:** <https://planetscale.com>

## API Keys (Get Free)

- **Groq:** <https://console.groq.com>
- **OpenAI:** <https://platform.openai.com>
- **Together AI:** <https://api.together.xyz>

## Learning Resources

- **Flask Tutorial:** <https://flask.palletsprojects.com/tutorial/>
- **Python Docs:** <https://docs.python.org/3/>
- **MDN Web Docs:** <https://developer.mozilla.org/>

---

## Next Steps

### For Beginners

1. Set up Ollama locally
2. Create simple Flask app
3. Build basic chat UI
4. Test locally
5. Deploy to Render (free)

## For Intermediate

1. Add user authentication
2. Implement chat history
3. Add file upload
4. Deploy with CI/CD
5. Monitor usage

## For Advanced

1. Add Redis caching
2. Implement WebSockets
3. Multi-model support
4. Advanced analytics
5. Scale with Kubernetes

---

## Summary Checklist

### Development Setup

- ☐ Install Python 3.9+
- ☐ Install Ollama
- ☐ Pull Ollama models
- ☐ Create virtual environment
- ☐ Install dependencies
- ☐ Create .env file
- ☐ Test locally

### Production Deployment

- ☐ Get Groq API key (free)
- ☐ Create Neon PostgreSQL database (free)
- ☐ Sign up for Render (free)
- ☐ Push code to GitHub
- ☐ Deploy to Render
- ☐ Add environment variables
- ☐ Test production deployment

## Post-Deployment

- ☐ Monitor usage
- ☐ Check error logs
- ☐ Test rate limiting
- ☐ Backup database
- ☐ Update documentation

---

**This guide contains everything discussed. Save it for offline reference!**

**Questions? Review the appropriate section above or refer to the official documentation links.**